Selalib Notes - DRAFT -

Edwin Chacon-Golcher

July, 2011

Abstract

This is a draft for a guide to the facilities offered by Selalib, the Semi-Lagrangian Library. It serves as a record for module descriptions, design decision, pending issues, etc., which are relevant for the development and improvement of the library. The main focus is the exposed API, thus the document also serves as a working document for the architectural specifications. More specific implementation details should be found in the comments within the source code. We may decide later to extract those either with an automatic system or in a developer's manual. At the time of writing this version of the draft, the library contains only a few modules in the lower levels and aims at having a minimalist structure. Of those capabilities that could be deemed interesting enough by the managers of GYSELA, the intent would be to incorporate them that code. In this sense, this library prototype can also be a testbed for ideas and suggestions that could eventually be absorbed into production code.

Contents

Chapter 1

Introduction

Selalib is the *Se*mi-*La*grangian *Lib*rary, a collection of types and its associated methods that are useful for creating parallel plasma physics simulations that use this specific methodology for the solution of the *Vlasov* equation. In its design, we have attempted to expose to the users an interface that expresses as naturally as possible the problems that arise when using the semi-lagrangian approach. The present version of the top-level types and interfaces is the same as discussed in Selalib's project meeting in January 2011.

Selalib is structured in layers. One can think about these layers as libraries in their own right. A given layer can use the capabilities offered by a lower layer through the exposed interface, but never from a higher level. The layer with the lowest level of abstraction presently contains basic utilities like memory allocators, assertions and basic numeric types. The second layer is composed of numerical and parallel utilities. The third and highest level of the library contains the semi-lagrangian methodology tools, types and methods. This manual will ultimately describe each of these layers.

Most of the native types and operations provided by Selalib are prefixed by `sll_`. In this way you can at least have an expectation of finding documentation (if a user) and a starting point of where to start looking if you wish to dive into some particular aspect of an implementation (if a developer). This is an early implementation, so the prefix to designate Selalib's features could change. It is also worth considering to eliminate the prefix for the lower layers of the library, as these are general/reusable components that could in principle be used in different projects. In such case, the `sll_` prefix would be reserved for the top-level layer that implements the specific semi-lagrangian functionality.

The Selalib prototype is written in Fortran 95. There are some aspects of its implementation that may warrant some commentary as these aspects have an impact on how the library is used. For instance, the full library would be imported by the declaration:

```
#include "selalib.h"
```

instead of the more Fortran-natural form:

```
use selalib
```

The reason for this is that some features of the library are implemented as *macros*. To the user of the library, it makes no difference whether some functionality is implemented in the form of a procedure or a macro, with the exception that presently, macro names are written in `ALL-CAPS` (but we could change this). The use of the macro is required to offer certain capabilities, like informative error messages. For a developer, the use of the macros is needed in many cases to reduce code redundancy. The need for the use of macros will hopefully be more understandable when the reader sees the behavior of calls to simple macros like `SLL_ALLOCATE()` or `SLL_ASSERT()`.

A macro is a pre-processor directive. Selalib uses only very simple macros that are handled by *fpp*, the Fortran Pre-Processor. *fpp*'s capabilities are very limited and one peculiarity of its output is that macros are expanded into a single long line, which can easily surpass the 132 character limit that Fortran systems have. For this reason alone, the compilation of Selalib requires the use of a compilation flag: `-ffree-line-length-none` (in **gfortran**) or its equivalent in another compiler. On a similar vein, some compilers require the extension `.F90` (as opposed to `.f90`) in order to apply a preprocessing step. Thus, all files in Selalib use the `.F90` extension, and this is another reason that clients of this version of the library should as well.

Chapter 2

Low-Level Layer: Basic Utilities

2.1 Numeric Types

2.1.1 Description

As a *convenience*, Selalib offers aliases to some of Fortran's basic numeric types. This is intended to:

1. concisely and uniformly make clear the intended precision of a given variable,

2. permit mixed precision representations when needed (for example, a developer could wish to represent a particular real number with a combination of a 32-bit integer and a 32-bit real instead of a single 64-bit real as is sometimes done in ultra-high performance software),

3. provide a centralized location to change the precision of a numerical representation program-wide, and

4. save from the typing of a few characters.

2.1.2 Exposed Interface

Selalib's numeric precision features are accessed through the following aliases:

alias	shorthand for...
`sll_int32`	`integer(kind=i32)`
`sll_int64`	`integer(kind=i64)`
`sll_real64`	`real(kind=f32)`
`sll_real64`	`real(kind=f64)`
`sll_int`	`integer(kind=user)`
`sll_real`	`real(kind=user)`
`sll_comp32`	`complex(kind=f32)`
`sll_comp64`	`complex(kind=f64)`

Where the kind type parameters `i32`, `i64`, `f32` and `f64` have been defined to give a representation that is at least the denoted size for a given number. `user` is available for a flexible kind type parameter. These kind type parameters can also be used to specify the precision of numerical constants in the usual Fortran way, i.e. `3.14159265_f32`.

2.1.3 Usage

To use the module in a stand-alone way, use the line:

```
#include "sll_working_precision.h"
```

The aliases are to be used like the native Fortran types that they are aliasing:

```
    sll_real64 :: my_pi
    sll_real64 :: theta
    sll_int32  :: i
    sll_int32  :: N
    sll_real64, dimension(:), allocatable :: w
    ! allocate w ...
    my_pi = 3.14159265358979932384626433_f64
    theta = 2.0*my_pi/real(N,f64)
    do i=1,N/2
        w(i) = exp((0.0,1.0_f64)*theta*real(i,f64))
    end do
```

2.1.4 Status

Unit-tested.

2.2 Memory Allocator

2.2.1 Description

Selalib's memory allocators are simple wrappers around Fortran's native allocators. We ask very little from these allocators:

1. to allocate memory,

2. to initialize it if requested (only for Fortran-native types),

3. to deallocate memory, and

4. to fail with as descriptive a message as possible.

2.2.2 Exposed Interface

The interface to these allocators follows very closely the interface of Fortran's `allocate()` and `deallocate()` functions. The user may decide to allocate arrays or pointers of any type, and up to the same number of dimensions as permitted by a given Fortran implementation. The exposed macros are:

```
SLL_ALLOCATE( array_and_lims, ierr )
```

This is the basic memory allocator with the same syntax as Fortran's native `allocate()` but with a required integer error parameter. Any type and dimension can be given as an argument.

```
SLL_CLEAR_ALLOCATE( array_and_lims, ierr )
```

Same behavior as `SLL_ALLOCATE()` but also initializes the allocated memory to zero. This works for Fortran native types only or for derived types for which an assignment operator ($= 0$) has been defined.

```
SLL_DEALLOCATE( array_and_lims, ierr )
```

Basic deallocator. It is a wrapper around Fortran's native `deallocate()` function but also nullifies the pointer given as an argument after deallocation.

```
SLL_DEALLOCATE_ARRAY( array, ierr )
```

The array deallocator differs from the previous in that it does not attempt to nullify the array name after the deallocation is complete. This and the previous macro could in principle be merged into one for simplification.

```
SLL_INIT_ARRAY( array, val )
```

While not really an allocator/deallocator, we expose this macro for convenience in initializing an array with a given constant value. For consistency, it may be decided to eliminate this from the interface.

In contrast with the native `allocate()`, Selalib's allocators require that the user provide an integer variable for error checking (`ierr`).

2.2.3 Usage

To use the memory module in stand-alone fashion, use the line:

```
#include "sll_memory.h"
```

An example of the use of the module is:

```
    integer :: err
    real, dimension(:), allocatable :: a
    real, dimension(:,:,:), pointer :: b=>null()

    SLL_ALLOCATE( a(5000), err )
    SLL_CLEAR_ALLOCATE( b(1:4,1:3,1:2), err)
    SLL_INIT_ARRAY(b,0)
```

When finding an error condition, the user is informed about the location of the failing call:

```
Memory allocation Failure. Triggered in FILE tester.F90,
in LINE:    35
STOP ERROR: test_error_code(): exiting program
```

2.2.4 Status

Unit-tested.

2.3 Assertions

2.3.1 Description

This is a very small but very useful capability that permits the developer to devise many sorts of defensive programming techniques. The simple idea is that of declaring a condition that is expected to be true, thus triggering a descriptive error if the condition is violated.

2.3.2 Exposed Interface

Wherever a specific condition needs to be asserted, simply write:

```
    SLL_ASSERT( logical_condition )
```

Thus the condition to be checked must be cast in the form of a logical statement. Falseness of such statement would trigger an assertion error.

The assertions can be used liberally since they are controlled by a -DEBUG flag at compilation time. Absence of this flag will delete all calls to SLL_ASSERT() from the code. Hence, assertion conditions can be used during a debugging or testing phase without increasing the overhead in production code.

2.3.3   Usage

To use the *assertions* module in a stand-alone way, use the line:

```
#include "sll_assert.h"
```

However, all the following steps are necessary:

1. to make sure that the macros expand properly when using **gfortran**, include the compilation flag `-DGFORTRAN`. This is to ensure that the assertion conditions can be properly converted into strings to be returned in case that an assertion is triggered,

2. when using **gfortran**, use pass along the flag `-ffree-line-length-none` to prevent compilation error as some of the included macros may go beyond Fortrans 132 character limit,

3. to activate the assertions, pass along the flag `-DDEBUG`. If this flag is not present, the assertions will compile into nothing, which may be desirable after the code has been debugged.

An example may be the checking of in-range indices on a protected array:

```
function get_val( a, i )
     sll_int32 :: get_val
     sll_int32, intent(in) :: i
     sll_int32, dimension(:), intent(in) :: a
     SLL_ASSERT( (i .ge. 1) .and. (i .le. size(a)) )
     get_val = a(i)
end function get_val
```

Which could produce a behavior such as:

```
$./unit_test
    The size of a is: 1000
    a(1) =    0
    a(117) =    0
    Here we ask for the value of a(1001):

    (i .ge. 1) .and. (i .le. size(a)) : Assertion
    error triggered in file unit_test.F90 in line       25
    STOP :  ASSERTION FAILURE
```

Such way of stopping a program is much less uncomfortable than the sinking feeling one has when the program stops as in:

```
$ ./unit_test
    Array values:
    The size of a is: 1000
    a(1) =    0
    a(117) =    0
    Incident de segmentation (core dumped)
```

9

which of course doesn't even need to happen at the moment of the first error.

2.3.4   Status

Unit tested.

Chapter 3

Mid-Level Layer: Numerical and Parallel Utilities

3.1 Tridiagonal System Solver

3.1.1 Description

To solve systems of the form $Ax = b$, where A is a tridiagonal matrix, Selalib offers a native, robust tridiagonal system solver. The present implementation contains only a serial version. The algorithm is based on an LU factorization of a given matrix, with row pivoting. The tridiagonal matrix must be given as a single array, with a memory layout shown next.

$$\begin{bmatrix} a(2) & a(3) & & & & & & a(1) \\ a(4) & a(5) & a(6) & & & & & \\ & a(7) & a(8) & a(9) & & & & \\ & & \ddots & \ddots & & \ddots & & \\ & & & \ddots & & \ddots & & \ddots \\ & & & & a(3n-5) & a(3n-4) & a(3n-3) \\ a(3n) & & & & & a(3n-2) & a(3n-1) \end{bmatrix}$$

3.1.2 Exposed Interface

Factorization of the matrix A is obtained through a call to the subroutine

```
sll_setup_cyclic_tridiag( a, n, lu, ipiv )
```

where `a` is the matrix to be factorized, `n` is the problem size (the number of unknowns), `lu` is a real array of size $7n$ where factorization information will be

returned and `ipiv` is an integer array of length `n` on which pivoting information will be returned. From the perspective of the user, `lu` and `ipiv` are only arrays that `sll_setup_cyclic_tridiag` requires and do not need any further consideration.

The solution of a tridiagonal system, once the original array A has been factorized, is obtained through a call to

```
sll_solve_cyclic_tridiag( lu, ipiv, b, n, x )
```

where `lu`, `ipiv` are the arrays returned by `sll_setup_cyclic_tridiag()`, `b` is the independent term in the original matrix equation, `n` is the system size and `x` is the array where the solution will be returned.

3.1.3   Usage

To use the module in a stand-alone way, include the line:

```
use sll_tridiagonal
```

The following code snippet is an example of the use of the tridiagonal solver.

```
sll_int32 :: n = 1000
sll_int32 :: ierr
sll_real64, allocatable, dimension(:) :: lu
sll_int32,  allocatable, dimension(:) :: ipiv
sll_real64, allocatable, dimension(:) :: x

SLL_ALLOCATE( lu(7*n), ierr )
SLL_ALLOCATE( ipiv(n), ierr )
SLL_ALLOCATE( x(n), ierr )

! initialize a(:) with the proper coefficients here...
and then:

sll_setup_cyclic_tridiag( a, n, lu, ipiv )
sll_solve_cyclic_tridiag( lu, ipiv, b, n, x )

 SLL_DEALLOCATE_ARRAY( lu, ierr )
 SLL_DEALLOCATE_ARRAY( ipiv, ierr )
 SLL_DEALLOCATE_ARRAY( x, ierr )
```

Note that if the last call had been made as in

```
sll_solve_cyclic_tridiag( lu, ipiv, b, n, b )
```

the system would have been solved in-place.

3.1.4   Status

Unit-tested.

3.2 Splines

3.2.1 Description

The splines module provides capabilities for 1D data interpolation with cubic B-splines and different boundary conditions (at the time of this writing: periodic, hermite). The data to be interpolated is represented by a simple array. The spline coefficients and other information are stored in a spline object, which is also used to interpolate the fitted data.

3.2.2 Exposed Interface

Fundamental type:

```
sll_spline_1d
```

Like all the other fundamental types in Selalib, this type is declared as a pointer and only manipulated through the functions and subroutines described below. For more explicit examples, see the usage section.

Available functions:

```
new_spline_1D( data, num_cells, xmin, xmax, bc_type )
```

`new_spline_1D()` is responsible for allocating all the necessary storage for the spline object and its initialization. Hence the full creation of a spline is a two-step process. The first one is the declaration of the spline pointer, the second is its assignment through a call to the `new_spline_1D()` function. For details, see the usage section below. In this function:

data: The double precision floating point data array to be fitted.

num_cells: The data array represents the values of a function at the nodes of a uniform 1D grid. `num_cells` represents the number of cells in this uniform grid. For periodic boundary conditions one may choose that the number of cells equals the number of data points in the array. In other words, for a data array with points indexed 1:NP, we can think that these data represent a grid with NP cells, and where the following point (NP+1) would be the same as point 1. In case of hermitian boundary conditions, a data array indexed 1:NP would correspond to a grid of size NP-1 cells, as both endpoints are included.

xmin: Lower bound of the grid in which the data array is defined. In other words, if we think of the data array (indexed 1:NP) as the values of a discrete function f defined over a sequence of x_i's, then $x_1 = xmin$.

xmax: Similarly to `xmin`, `xmax` represents the maximum extent of the grid. Notice, however, that this is influenced by how one thinks of the data in the grid; for example, consider the case of periodic data defined over a grid of NC cells and indexed 1:NC (that is, the point NC+1 is the same as

the point 1 and is not taken into account). In this case, $xmax = x_{NC+1}$. In such way that $xmin - xmax$ is really the period of the function. In the case of data indexed 1:NC+1, and hermitian boundary conditions, then $xmax = x_{NC+1}$, the location of the last point.

bc_type: Descriptor of the type of boundary condition to be imposed in the spline. Presently, one of `PERIODIC_SPLINE` or `HERMITE_SPLINE`. These are really aliases to integer flags, but here we avoid the use of non-descriptive flags.

```
interpolate_value( x, spline )
```

`interpolate_value()` returns the value of *f(x)* where x is a floating-point value between `xmin` and `xmax`, and f is a continuous function built with cubic B-splines and user-defined boundary conditions. Essentially, the spline object and the interpolate function create the illusion of having available a continuous function when originally, only discrete data were available. Available subroutines:

```
compute_spline_1D( data, num_cells, bc_type, spline )
```

`compute_spline_1D()` is responsible for actually computing the spline information (i.e.: the spline coefficients) when one does not wish to allocate a new object. This can be thought of a way to *update* the spline coefficients of a given spline object. In this subroutine:

data: The double precision floating point data array to be fitted.

num_cells: The number of cells of the grid that underlies the data array.

bc_type: Descriptor of the type of boundary condition to be imposed in the spline. Presently, one of `PERIODIC_SPLINE` or `HERMITE_SPLINE`. These are really aliases to integer flags, but here we avoid the use of non-descriptive flags.

spline: A pointer to the spline object to be filled or updated.

```
compute_spline_1D_periodic( data, num_cells, spline )
compute_spline_1D_hermite ( data, num_cells, spline )
```

These are subroutines are meant as an alternative interface to the services of `compute_spline_1D()`, in which the required boundary condition is explicit in the name of the function. These functions are slightly faster in that the programmer is choosing the boundary condition, and thus a branch can be avoided in the former function.

```
delete_spline_1D( spline )
```

`delete_spline_1D()` is the spline destructor. Its use is important to avoid memory leaks. As usual, dummy arguments that refer to fundamental types (like spline) are pointers to such types.

3.2.3 Usage

To use the module in a stand-alone way, include the line:

```
use sll_splines
```

The following example is an extract from the module's unit test.

```
1      program spline_tester
2      #include "sll_working_precision.h"
3      #include "sll_assert.h"
4      #include "sll_memory.h"
5        use sll_splines
6        use numeric_constants
7        implicit none
8
9      #define NC 32
10
11       sll_int32 :: err
12       sll_int32 :: i
13       type(sll_spline_1d), pointer :: sp1
14       type(sll_spline_1d), pointer :: sp2
15       sll_real64, allocatable, dimension(:) :: data
16       sll_real64 :: accumulator1, accumulator2
17
18       accumulator1 = 0.0_f64
19       accumulator2 = 0.0_f64
20
21       SLL_ALLOCATE(data(NC+1), err)
22
23       print *, 'initialize data array'
24       do i=1,NC+1
25         data(i) = sin((i-1)*sll_pi/real(NC,f64))
26       end do
27
28       sp1 =>  new_spline_1D( data(1:NC), NC, 0.0_f64, sll_pi, PERIODIC_SPLINE )
29       sp2 =>  new_spline_1D( data(1:(NC+1)), NC, 0.0_f64, sll_pi, HERMITE_SPLINE )
30
31       print *, 'cumulative errors at nodes: '
32       do i=1, NC+1
33          accumulator1 = accumulator1 + abs(data(i) - &
34               interpolate_value(real(i-1,f64)*sll_pi/real(NC,f64), sp1))
35       end do
36
37       print *, 'hermite case: '
38       do i=1, NC+1
39          accumulator2 = accumulator2 + abs(data(i) - &
```

```
40                 interpolate_value(real(i-1,f64)*sll_pi/real(NC,f64), sp2))
41        end do
42        print *, 'Periodic case: '
43        print *, 'average error at the nodes = '
44        print *, accumulator1/real(NC,f64)
45        call delete_spline_1D(sp1)
46        if( accumulator1/real(NC,f64) < 1.0e-15 ) then
47           print *, 'PASSED TEST'
48        else
49           print *, 'FAILED TEST'
50        end if
51        print *, '**************************** '
52        print *, 'Hermite case: '
53        print *, 'average error at the nodes = '
54        print *, accumulator2/real(NC,f64)
55        call delete_spline_1D(sp2)
56        if( accumulator2/real(NC,f64) < 1.0e-15 ) then
57           print *, 'PASSED TEST'
58        else
59           print *, 'FAILED TEST'
60        end if
61     end program spline_tester
```

Here we do not go in detail over every line but only highlight those lines in which we interact with the splines module

Line 5: Imports the spline module. The intent is to eventually not require this but to import a single module, say 'selalib' which will include all modules itself. For now, this is the way to include these individual capabilities.

Lines 13 - 14: Declaration of the spline pointers.

Lines 28 - 29: Allocation and initialization of the splines. Note the `=>` pointer assignment syntax.

Lines 34, 40: Value interpolation using the existing splines.

Lines 45, 55: Destruction of spline objects.

In this example we do not show the use of the `compute_spline_1D()` functions. These are used though, albeit indirectly. For example, the initialization of a spline object uses the `compute_spline_1D()` internally. We could have also at some point in the program made a direct call such as

```
call compute_spline_1D( other_data, NC, PERIODIC_SPLINE, sp1 )
```

and updated the contents of `sp1`.

3.2.4 Status

Unit-tested.

3.3 Gauss-Legendre Integrator

3.3.1 Description

This is a low-level mathematical utility that applies the Gauss-Legendre method to compute numeric integrals. This module aims at providing a single interface to the process of integrating a function on a given interval.

3.3.2 Exposed Interface

To integrate the function `f(x)` (real-valued and of a single, real-valued argument `x`) over the interval $[a, b]$, the simplest way is through a function call such as:

```
gauss_legendre_integrate_1D(f, a, b, n)
```

In the function above, `n` represents the desired number of *Gauss* points used in the calculation:

$$\int_{-1}^{1} f(x)\mathrm{d}x \approx \sum_{k=1}^{n} w_k f(x_k) \tag{3.1}$$

Presently, the implementation accepts values of `degree` between 2 and 10 inclusively. The function `gauss_legendre_integrate_1D` internally does the proper scaling of the points to adjust the integral over the desired interval.

The function `gauss_legendre_integrate_1D` is a generic interface and as such, it hides some alternative integrators, which are selected depending on the type of the passed arguments. For instance, we have available the function

```
gauss_legendre_integrate_interpolated_1D(f, spline, a, b, n)
```

which integrates a function represented by a spline object. The function `f` in this case is the spline interpolation function. It looks like this interface could be simplified and we could eliminate the first parameter and pass only the spline object. The only reason to leave the interpolation function as an argument is if we find some compelling reason to parametrize the interpolation function as well.

It will be necessary to implement other integrators for functions with a different signature, such as two- or three-parameter functions. It might also be necessary to distinguish between one-dimensional and two- or 3-dimensional integration. While this is not yet implemented, here we lay out some suggestions on how to proceed in such cases.

For the class of integrals that are done in one-dimension, such as the above, the cleanest but somewhat more laborious approach appears to be to write a different integrator for every function signature that is needed. For instance,

one may need to write specialized integrators for $f(x_1, x_2)$, $f(x_1, x_2, x_3)$ and so on, with the convention that the integral is carried out over, say, the first of the variables. The variables which are not integrated can be used as parameters. The alternative to this approach could be to write a single integrator that is able to receive multiple parameters, through the use of arrays or derived types, but the ugliness of this approach, and the need to basically write glue-code (pack/unpack the arrays or derived types with variables and parameters) every time one wants to integrate something are reasons to reject this approach.

3.3.3 Usage

As mentioned above, the name of the generic function that hides the specialized functions is `gauss_legendre_integrate_1D`. The specialized functions can be individually called to avoid the overhead of the generic function call if desired. A one-dimensional function (user or Fortran) can be integrated by a call like:

```
gauss_legendre_integral_1D( test_function, &
                            0.0_f64,        &
                            sll_pi/2.0,     &
                            4 )
```

A function that is represented by an underlying spline object can be called like:

```
gauss_legendre_integral_interpolated_1D( interpolate_value,&
                                         sp1,              &
                                         0.0_f64,          &
                                         sll_pi,           &
                                         4)
```

where sp1 is a spline object. It should be decided if this last case is indeed that interface that is wished, or if something more simplified should be implemented intstead.

3.3.4 Status

Unit-tested.

3.4 Collective Communications

3.4.1 Description

Selalib applies the principle of modularization throughout all levels of abstraction of the library and aims at keeping third-party library modules as what they are: separate library modules. Therefore, in its current design, even a library like MPI has a single point of entry to Selalib. The collective communications module is such point of entry. We focus thus on the functionality offered by MPI, assign wrappers to its most desirable functionalities and write wrappers around

them. These are the functions that are actually used throughout the program. This allows to adjust the exposed interfaces, do additional error-checking and would even permit to completely change the means to parallelize a code, by being able to replace MPI in a single file if this were ever needed.

3.4.2   Exposed Interface

Fundamental type:

```
sll_collective_t
```

Constructors, destructors and access functions:

```
sll_new_collective( parent_col )
sll_delete_collective( col )
```

When the Selalib environment is activated, there exists, in exact analogy with `MPI_COMM_WORLD`, a global named `sll_world_collective`. At the beginning of a program execution, this is the only collective in existence. Further collectives can be created down the road. The above functions are responsible for the creation and destruction of such collectives. The following functions are used to access the values that a particular collective knows about.

```
sll_get_collective_rank( col )
sll_get_collective_size( col )
sll_get_collective_color( col )
sll_get_collective_comm( col )
```

Since the wrapped library requires initialization, so does `sll_collective`. To start and end the parallel environment, the user needs to call the functions:

```
sll_boot_collective( )
sll_halt_collective( )
```

These functions would not be exposed at the top level, and would be hidden by a further call to something akin to `boot_selalib` and `halt_selalib`. Finally, the wrappers around the standard `MPI` capabilities are presently exposed through the following generic functions:

```
sll_collective_bcast( col, buffer, size, root )
sll_collective_gather( col, send_buf, send_sz, root,
                       rec_buf )
sll_collective_gatherv( col, send_buf, send_sz, recvcnts,
                        displs, root, recv_buf )
sll_collective_allgatherv( col, send_buf, send_sz, sizes,
                           displs, rec_buf )
sll_collective_scatter( col, send_buf, size, root,
                        rec_buf )
sll_collective_scatterv( col, send_buf, sizes, displs,
```

```
                              rec_szs, root, rec_buf )
      sll_collective_all_reduce( col, send_buf, count, op,
                              rec_buf )
```

which presently stand for specialized versions that operate on specific types.
For instance:

```
      sll_collective_all_to_allv_real( send_buf,
                                     send_cnts,
                                     send_displs,
                                     recv_buf,
                                     recv_cnts,
                                     recv_displs,
                                     col )
```

3.4.3   Usage

To use the module as stand-alone, include the line:

```
      use sll_collective
```

Any use of the module's functionalities must be preceeded by calling

```
      call sll_boot_collective()
```

and to "turn off" the parallel capabilities, one should finish by a call to:

```
      call sll_halt_collective()
```

This *booting* of the parallel environment needs to be done only once in a program.
   Some more specific examples are needed here...

3.4.4   Status

Several core functionalities tested, but no comprehensive unit test done yet

3.5   Remapper

3.5.1   Description

Written on top of `sll_collective`, the remapper is a powerful facility that is
capable of rearranging data in flexible and convenient ways in a parallel machine.
It is meant to be a generalization of the 'transposition', which users/developers
of *CALVI* team codes know and love. The main difference is is generality, as
here we extend the idea to encompass something beyond a data transposition in
2D, to an operation that can be carried out in any number of dimensions. For
instance, suppose that you start with a multidimensional array that has been
domain decomposed and distributed among N_p processors. The layout of the
data (that is, the description of what ranges of the data are contained in each

processor) is specified by an instance of the type `layout_XD_t`, (where `X` is the dimension of the data). The layout contains a notion of an N_p-sized collection of boxes, each box representing a contiguous chunk of the multidimensional array stored in each node. If in the course of a computation, you wish to reconfigure the layout of the data (for example, if you wished to re-arrange data in a way that would permit launching serial algorithms locally in each node), then you would create and initialize a new layout descriptor with the target configuration (i.e.: you to define the box to be stored in each node). This is a conceptually simple but perhaps slightly verbose task. Then a call to the appropriate choice among:

```
NEW_REMAPPER_PLAN_3D( initial_layout,
                      target_layout,
                      data_size_in_integer_sizes )
NEW_REMAPPER_PLAN_4D( initial_layout,
                      target_layout,
                      data_size_in_integer_sizes )
NEW_REMAPPER_PLAN_5D( initial_layout,
                      target_layout,
                      data_size_in_integer_sizes )
```

will yield an instance of the type `remap_plan_3D_t`, or `remap_plan_4D_t` or `remap_plan_5D` , respectively, that will contain all the information necessary to actually carry out the data re-distribution. Finally, a call to

```
apply_remap_3D( plan, data_in, data_out )
apply_remap_4D( plan, data_in, data_out )
apply_remap_5D( plan, data_in, data_out )
```

will actually redistribute `data` (as an out-of-place operation) according to `plan` in an optimized way[1].

To appreciate the power of such facility, note that in principle, the construction of a (communications latency-limited) parallel quasi-neutral solver can be based exclusively on remapping operations. This is an important tool in any problem that would require global rearrangements of data. The remapper thus is able to present a single powerful abstraction that is general, reusable and completely hides most of the complications introduced by the data distribution.

3.5.2   Exposed Interface

The remapper offers the following descriptor types for parallel data layout, differing from one another only in the dimensionality of the data described:

---

[1]This is a very loaded comment. Some of the optimizations are carried out by the remapper, like the identification of the minimally-sized communicators to launch the exchanges, or the selection of the lower-level communications functions (alltoall vs. alltoallv, for instance). Other optimizations would need to be triggered externally, by passing proper compilation flags to the MPI facilities. This would be problem-dependent.

```
layout_3D_t
layout_4D_t
layout_5D_t
```

(Note that for the remapper, we have forgone the use of the `sll_` prefix. This is as an example of the likely policy that the low- and mid-level reusable utilities should not be prefixed, thus being instantly available for any other development. Eventually a decision needs to be made and the choice implemented uniformly throughout the library.) These types are each accompanied by their own constructors, destructors and accessors. Specifically, the constructors are:

```
new_layout_3D( collective )
new_layout_4D( collective )
new_layout_5D( collective )
```

Note that each layout descriptor needs to be allocated by providing an instance of `sll_collective_t`. This can be understood by thinking of the data layout as being associated with a given group of processors (the collective) and a specification of the data boxes contained in each one. After calling any of the `new_layout` functions, the returned instance becomes associated to the given collective and enough memory is allocated (size of the collective) to hold the boxes specification.

The destructors are:

```
delete_layout_3D( layout )
delete_layout_4D( layout )
delete_layout_5D( layout )
```

The access functions for the `layout` types are are always prefixed with the corresponding `get_layout_XD`/`set_layout_XD` (where the 'X' denotes the dimensionality of the data), and they presuppose knowledge of the convention for ordering the indices as in `i, j, k, l, m`, for the dimensions. Specifically, to get/set values inside the `layout` types we have available for 3D layouts:

```
get_layout_3D_num_nodes( layout )
get_layout_3D_box( layout, rank )

get_layout_3D_i_min( layout, rank )
get_layout_3D_i_max( layout, rank )
get_layout_3D_j_min( layout, rank )
get_layout_3D_j_max( layout, rank )
get_layout_3D_k_min( layout, rank )
get_layout_3D_k_max( layout, rank )

set_layout_3D_i_min( layout, rank, val )
set_layout_3D_i_max( layout, rank, val )
set_layout_3D_j_min( layout, rank, val )
```

```
set_layout_3D_j_max( layout, rank, val )
set_layout_3D_k_min( layout, rank, val )
set_layout_3D_k_max( layout, rank, val )
```

As a very inelegant convenience, the layout type allows direct access to its collective reference. For 4D layouts:

```
get_layout_4D_num_nodes( layout )
get_layout_4D_box( layout, rank )

get_layout_4D_i_min( layout, rank )
get_layout_4D_i_max( layout, rank )
get_layout_4D_j_min( layout, rank )
get_layout_4D_j_max( layout, rank )
get_layout_4D_k_min( layout, rank )
get_layout_4D_k_max( layout, rank )
get_layout_4D_l_min( layout, rank )
get_layout_4D_l_max( layout, rank )

set_layout_4D_i_min( layout, rank, val )
set_layout_4D_i_max( layout, rank, val )
set_layout_4D_j_min( layout, rank, val )
set_layout_4D_j_max( layout, rank, val )
set_layout_4D_k_min( layout, rank, val )
set_layout_4D_k_max( layout, rank, val )
set_layout_4D_l_min( layout, rank, val )
set_layout_4D_l_max( layout, rank, val )
```

And for 5D layouts:

```
get_layout_5D_num_nodes( layout )
get_layout_5D_box( layout, rank )

get_layout_5D_i_min( layout, rank )
get_layout_5D_i_max( layout, rank )
get_layout_5D_j_min( layout, rank )
get_layout_5D_j_max( layout, rank )
get_layout_5D_k_min( layout, rank )
get_layout_5D_k_max( layout, rank )
get_layout_5D_l_min( layout, rank )
get_layout_5D_l_max( layout, rank )
get_layout_5D_m_min( layout, rank )
get_layout_5D_m_max( layout, rank )

set_layout_5D_i_min( layout, rank, val )
set_layout_5D_i_max( layout, rank, val )
set_layout_5D_j_min( layout, rank, val )
```

```
set_layout_5D_j_max( layout, rank, val )
set_layout_5D_k_min( layout, rank, val )
set_layout_5D_k_max( layout, rank, val )
set_layout_5D_l_min( layout, rank, val )
set_layout_5D_l_max( layout, rank, val )
set_layout_5D_m_min( layout, rank, val )
set_layout_5D_m_max( layout, rank, val )
```

The above functions define the interface that will allow you to declare and initialize the `layout` types as desired. This is where the work lies when using this module. Note that all the above functions could be coalesced into a set of functions of the type `set_layout_X_XXX(layout, rank, val)` if we choose to hide all the above functions behind a generic interface. The selection would be done automatically depending on the type of layout passed as an argument.

The type `remap_plan` exists also in multiple flavors, depending on the dimensionality of the data to be remapped:

```
remap_plan_3D_t
remap_plan_4D_t
remap_plan_5D_t
```

The `remap_plan_t` type stores the locations of the memory buffers that will be involved in the communications, the specification of the data that will be sent and received, as well as the collective within which the communications will take place. There are, however, declaration functions available. The choice depends on the dimensionality of the data:

```
NEW_REMAPPER_PLAN_3D( initial_layout,
                      final_layout,
                      array_name )
NEW_REMAPPER_PLAN_4D( initial_layout,
                      final_layout,
                      array_name )
NEW_REMAPPER_PLAN_5D( initial_layout,
                      final_layout,
                      array_name )
```

Finally, the way to execute the plan on a particular data set is through a call of the appropriate subroutine (here presented as generic interfaces)

```
apply_remap_3D( plan, data_in, data_out )
apply_remap_4D( plan, data_in, data_out )
apply_remap_5D( plan, data_in, data_out )
```

3.5.3   Usage

For use in stand-alone way, use the line:

```
#include "sll_remap.h"
```

While verbose, the best way to demonstrate the usage of the remapper is with
a complete program. Below it, we examine the different statements.

```
1    program remap_test
2      use sll_collective
3    #include "sll_remap.h"
4    #include "sll_memory.h"
5    #include "sll_working_precision.h"
6    #include "misc_utils.h"
7      implicit none
8
9      ! Test of the 3D remapper takes a 3D array whose global
10     ! size Nx*Ny*Nz, distributed among pi*pj*pk processors.
11     integer, dimension(:,:,:), allocatable :: a3
12     integer, dimension(:,:,:), allocatable :: b3
13
14     ! Take a 3D array of dimensions 8X8X1
15     integer, parameter                :: total_sz_i = 8
16     integer, parameter                :: total_sz_j = 8
17     integer, parameter                :: total_sz_k = 1
18
19     ! the process mesh
20     integer, parameter                :: pi = 4
21     integer, parameter                :: pj = 4
22     integer, parameter                :: pk = 1
23
24     ! Split it in  16 processes, each with a local chunk 2X2X1
25     integer                           :: local_sz_i
26     integer                           :: local_sz_j
27     integer                           :: local_sz_k
28     integer                           :: ierr
29     integer                           :: myrank
30     integer                           :: colsz  ! collective size
31     integer                           :: i,j,k
32     integer                           :: i_min, i_max
33     integer                           :: j_min, j_max
34     integer                           :: k_min, k_max
35     integer                           :: node
36     integer, dimension(1:3)           :: gcoords
37
38     ! Remap variables
39     type(layout_3D_t), pointer        :: conf3_init
40     type(layout_3D_t), pointer        :: conf3_final
41     type(remap_plan_3D_t), pointer    :: rmp3
42
```

```
43    ! Boot parallel layer
44    call sll_boot_collective()
45
46    ! Initialize and allocate the variables.
47    local_sz_i = total_sz_i/pi
48    local_sz_j = total_sz_j/pj
49    local_sz_k = total_sz_k/pk
50    SLL_ALLOCATE(a3(1:local_sz_i,1:local_sz_j,1:local_sz_k), ierr)
51    SLL_ALLOCATE(b3(1:local_sz_i,1:local_sz_j,1:local_sz_k), ierr)
52    myrank    = sll_get_collective_rank(sll_world_collective)
53    colsz     = sll_get_collective_size(sll_world_collective)
54
55    conf3_init    => new_layout_3D( sll_world_collective )
56    conf3_final   => new_layout_3D( sll_world_collective )
57    random_layout1 => new_layout_3D( sll_world_collective )
58
59    ! Initialize the layout
60    do k=0, pk-1
61       do j=0, pj-1
62          do i=0, pi-1
63             node = i+pi*(j+pj*k) ! linear index of node
64             i_min = i*local_sz_i + 1
65             i_max = i*local_sz_i + local_sz_i
66             j_min = j*local_sz_j + 1
67             j_max = j*local_sz_j + local_sz_j
68             k_min = k*local_sz_k + 1
69             k_max = k*local_sz_k + local_sz_k
70             call set_layout_i_min( conf3_init, node, i_min )
71             call set_layout_i_max( conf3_init, node, i_max )
72             call set_layout_j_min( conf3_init, node, j_min )
73             call set_layout_j_max( conf3_init, node, j_max )
74             call set_layout_k_min( conf3_init, node, k_min )
75             call set_layout_k_max( conf3_init, node, k_max )
76          end do
77       end do
78    end do
79
80    ! Initialize the data. We use the information in the layout.
81    do k=1, local_sz_k
82       do j=1, local_sz_j
83          do i=1, local_sz_i
84             gcoords =  local_to_global_3D( conf3_init, (/i,j,k/) )
85             a3(i,j,k) = gcoords(1) + &
86                   total_sz_i*((gcoords(2)-1) + &
87                   total_sz_j*(gcoords(3)-1))
88          end do
```

26

```
89        end do
90     end do
91
92     ! Initialize the final layout, in this case, just a
93     ! transposition
94     do k=0, pk-1
95        do j=0, pj-1
96           do i=0, pi-1
97              node = i+pi*(j+pj*k) ! linear index of node
98              i_min = i*local_sz_i + 1
99              i_max = i*local_sz_i + local_sz_i
100             j_min = j*local_sz_j + 1
101             j_max = j*local_sz_j + local_sz_j
102             k_min = k*local_sz_k + 1
103             k_max = k*local_sz_k + local_sz_k
104             call set_layout_i_min( conf3_final, node, j_min )
105             call set_layout_i_max( conf3_final, node, j_max )
106             call set_layout_j_min( conf3_final, node, i_min )
107             call set_layout_j_max( conf3_final, node, i_max )
108             call set_layout_k_min( conf3_final, node, k_min )
109             call set_layout_k_max( conf3_final, node, k_max )
110          end do
111       end do
112    end do
113
114    rmp3 => NEW_REMAPPER_PLAN_3D( conf3_init, conf3_final, a3 )
115    call apply_remap_3D( rmp3, a3, b3 )
116
117    ! At this moment, b3 contains the expected output from the &
118    ! remap operation.
119
120    ! Delete the layouts
121    call delete_layout_3D( conf3_init )
122    call delete_layout_3D( conf3_final )
123
124     call sll_halt_collective()
125
126  end program remap_test
```

Lines 1 - 5: Required preamble at the time of this writing. Eventually this will be replaced by a single statement to include the whole library. Presently, we include various headers individually, so bear in mind that this is not the way this will end up being. Line 3 specifically loads the remapper facility. Here it is brought as a header file as the NEW_REMAPPER_PLAN_XD() is implemented as a macro.

Lines 9 - 12: For this example we allocate two 3D arrays for the input and output of the remap operation.

Lines 14 - 22: Definition of the array size from a global perspective. In other words, the array to be remapped is a $8 * 8 * 1$ array, to be distributed on a processor mesh of dimensions $4 * 4 * 1$.

Lines 24 - 36 Miscellaneous integer variables that we will use.

Lines 38 - 41 Pointers to the initial and final layouts and the remap plan.

Line 44 Presently we boot from collective. Eventually this will be replaced by a call to something like `boot_selalib()` or something similar, where we declare and initialize anything we need in a single call.

Lines 46 - 57 Initialization of the variables.

Lines 59 - 78 This is where the actual work is when using the remapper. We need to initialize a layout, in this case the initial configuration. We use the access functions `set_layout_x_xxx( )` to populate the fields. Here we obviously take into account the geometry of our 'process mesh' to find out the rank of the process that we are initializing.

Lines 80 - 90 We need to initialize the data, here we choose simply to assign the index of the array, considered as a 1D array. Note the use of the helper function `local_to_global_3D( layout, triplet )`. We exploit the knowledge of the global layout of the data to find out the global indices of a local 3-tuple.

Lines 92 - 112 The other main part of the work, the initialization of the target layout. In this case, we chose a simple transposition, which is achieved by switching `i` and `j`.

Lines 114 - 115 Here we allocate and initialize the remap plan, using the initial and final configurations as input. The third argument is passed to inform the remapper of the type of data to be passed. The call to `apply_remap_3D()` is a call to a generic function, hence, a type-dependent sub-function must have been defined to be able to successfully make this call. At the time of this writing, only single precision integers and double precision floats have been implemented.

Line 116 Here we apply the plan. This function is type-dependent due to the input/output arrays. Please refer to the implementation notes for some commentary on our options with this interface.

Lines 122 - 125 Cleanup. The layouts need to be deleted to prevent memory leaks.

3.5.4   Implementation Notes

The biggest challenge with the remapper is to attain a desired level of genericity and to preserve the modularity of the library. These two problems are intimately related. Ideally, we should be able to apply a remap operation on data of any type, including user-derived types. Another requirement has been to confine a library like MPI to a single entry point into Selalib. This means that we do not want the MPI derived types to pollute the higher abstraction levels of the library: especially at the top level, we want to express our programs with the capabilities of the Fortran language alone.

These requirements were solved in the prototype version of the remapper through the use of a single datatype to represent all other types of data at the moment of assembling the exchange buffers and launching the MPI calls. In our case, we have chosen to represent all data as 'integers'. This means that the exchange buffers that are stored in the remap plans are integer arrays. Thus, the design decision in the prototype has been to choose flexibility and ease of change over execution speed. In contrast with the C language, the constant call to the `transfer()` function to store and retrieve data from the exchange buffers carries with it a possibly significant execution time penalty.

The function `NEW_REMAPPER_PLAN_XD()` is by nature type-independent, as the design of the plan only depends on the layouts. However, it is also convenient to store the send/receive buffers in the plan, and the allocation of these buffers requires knowledge of the amount of memory required. This information is passed in the third argument. The macro will internally select an element of this array and determine its size in terms of the fundamental datatype being exchanged (i.e.: `integer`). This way we now how much memory to allocate in the buffers.

Another means to achieve the illusion of genericity are Fortran's built-in features in this regard. For example, we can have specialized `apply_remap_3D()` functions for the most commonly used datatypes, all hidden behind the same generic name. These specialized functions would not depend on the current choice of using a single type for the exchange buffers, eliminating any penalty that we are definitively paying at present, with the calls to the `transfer()` intrinsic function. This solution would mean writing redundant code, something that could be addressed with preprocessor macros, but this would not be a solution for eliminating the penalizations of the `transfer()` intrinsic when we are exchanging derived types. A solution that can exchange these arbitrary data while not requiring the use of the MPI derived types at the higher levels is yet to be found. It could be that the Fortran way to solve this problem would be to accept the invasion of MPI at the higher levels...

3.5.5   Status

In testing.

Chapter 4

Top-Level Layer: Semi-Lagrangian Toolbox

4.1 Quasi-Neutral Equation Solver

4.1.1 Description

4.1.2 Exposed Interface

Fundamental type: None. It is a function that operates on other top-level types.
Function:

```
sll_solve_quasi_neutral_equation( electron_T_profile_3D,
                                  electron_n_profile_3D,
                                  charge_density,
                                  phi )
```

4.1.3 Usage

4.1.4 Status

4.2 Particle Distribution Function

4.2.1 Description

4.2.2 Exposed Interface

Fundamental type:

```
sll_distribution_function_t
```

All the fundamental types in the library are implemented as pointers. This choice has been made to ease the addition of Python bindings, in case that an even higher-level interface is desired some day.

Constructor and destructor:

```
sll_new_distribution_function( nr, ntheta, nphi, nvpar, mu )
sll_delete_distribution_function( f )
```

The constructor essentially limits itself to allocating the memory for the type. An initialization step is required afterwards:

```
sll_initialize_df( boundary_type_r,
                   boundary_type_vpar,
                   species_charge )
```

The accessors that get/set a particular value of a distribution function on a node are defined as macros to be able to keep a single interface while not risking a penalization when used in critical loops. Other queries on this type are implemented as ordinary functions. (Need to define this better, for instance, if some query functions would operate on integer arguments and/or real coordinates.)

```
SLL_GET_DF_VAL( i, j, k, l, df )
SLL_SET_DF_VAL( val, i, j, k, l, df )
sll_interpolate_df( r, theta, phi, vpar, mu )
sll_compute_derivative( f, r, theta, phi, vpar, mu )
sll_get_df_nr( df )
sll_get_df_nphi( df )
sll_get_df_ntheta( df )
sll_get_df_nvpar( df )
sll_get_df_mu( df )
```

The type also offers the services:

```
sll_compute_moments( df, ... )
```

4.2.3 Usage

4.2.4 Status

4.3 Advection Field

4.3.1 Description

4.3.2 Exposed Interface

Fundamental type:

```
sll_advection_field_3D_t
```

This implies that one of the options is to have multiple representations, for 3D, 2D, 1D.

4.3.3   Usage

4.3.4   Status

4.4   Advection

4.4.1   Description

4.4.2   Exposed Interface

Fundamental type: None. This is a function that operates on multiple top-level
types. Function:

```
sll_advect( distribution_function,
            advection_field,
            dt,
            space_mesh
            scheme )
```

Above, `scheme` is the functional parametrization of the various methods in use
(PSM, BSL, ...) and for which we need a standardized interface. The above
assumes that we can devise a standard functional interface.

4.4.3   Usage

4.4.4   Status